

Question #	RFQ Reference	Question	Government Response
1	Section L / Tech Spec 5.1	The solicitation requires GPU availability by May 15, 2026. Will the Government consider proposals where the infrastructure subcontractor can provision and deliver all 50 GPUs within 10 business days of contract award, assuming award occurs no later than May 5, 2026?	No. The Government will not consider proposals that require a 10-business-day provisioning period following contract award. The requirement remains that all 50 GPUs must be available by May 15, 2026, or immediately upon contract award if the award occurs after that date.
2	Section L, Vol I, Go/No-Go #3	The Go/No-Go requirement lists the NVIDIA HGX H100 80GB SXM5 as the baseline but accepts "equivalent or better" including Blackwell B200/GB200. Will the Government confirm that a proposal offering NVIDIA Blackwell B200 GPUs as the primary solution will receive a "Pass" on the Go/No-Go compliance check, given that B200 represents a superior architecture to the H100?	Yes, the Government confirms that a proposal offering NVIDIA Blackwell B200 GPUs satisfies the hardware baseline criteria. In accordance with Section L, Vol I, Go/No-Go #3, the Government will accept the NVIDIA HGX H100 80GB SXM5, an 'equivalent or better' architecture (including the B200), or older GPUs that are proven to perform specific applications as fast as newer GPUs. A final 'Pass' rating remains contingent upon the proposal demonstrating compliance with all stated technical and performance requirements.

3	Tech Spec 5.1	The 100 additional reserve GPUs must be "identical" to the primary 50. If the primary offer is Blackwell B200, must the reserve also be B200, or may the reserve be H100 (the stated baseline equivalent)?	The reserve GPUs may be H100, equivalent, or better; they do not strictly need to be identical to the primary 50 GPUs if the primary GPUs exceed the baseline. The Government has revised Go/No-Go Requirement 3: GPU Hardware & Capability in the solicitation for clarity to reflect this update.
4	Section L, Vol II	The Price Quote requires a single all-inclusive per-GPU-hour rate. Please confirm the exact GPU-hour quantity the Government will use to evaluate pricing for each CLIN. Specifically: - CLIN 0001 (May 15 – Oct 14, 2026): Is the evaluated quantity 183,600 GPU-hours (50 GPU Nodes x 24 hrs x 153 calendar days)? - CLIN 0002 (Oct 15, 2026 – Mar 14, 2027): Is the evaluated quantity 181,200 GPU-hours (50 GPU Nodes x 24 hrs x 151 calendar days)?	Yes, the Government confirms that for evaluation purposes only, the quantities stated are correct. The evaluated quantity for CLIN 0001 is 183,600 GPU-hours, and the evaluated quantity for CLIN 0002 is 181,200 GPU-hours. This has been clarified in the revised Section L, Factor 3 - Price.
5	Section M, Factor 3	Will the Government confirm the total evaluated price calculation is: CLIN 0001 GPU-hours + CLIN 0002 GPU-hours x per-GPU-hour rate? And that the two CLINs are evaluated separately based on actual calendar days in each period?	See answer to question 4.

6	Tech Spec 5.3 / Section L	The SOC 2 Type II certification is a Go/No-Go requirement. If the SOC 2 Type II certification is held by the subcontractor who operates the GPU infrastructure, will the Government accept the subcontractor's SOC 2 Type II certification as satisfying this Go/No-Go requirement?	Yes, provided the subcontractor's SOC 2 Type II certification covers the entirety of the infrastructure where Government data is hosted, processed, and stored. The proposal must make it clear that the prime contractor's systems will not host, process, or store any DARPA data.
7	Section L, Factor 2	May a prime contractor submit past performance references from its subcontractor, given that the subcontractor is the entity that previously performed GPU/HPC infrastructure contracts of similar scope and size?	Yes. The Government will accept past performance references from a prime contractor's proposed subcontractor
8	Tech Spec 5.2	Please clarify the expected scope of Tier 3 scientific software support. Specifically, does this include ML framework-level support (PyTorch, TensorFlow, JAX), or is it limited to infrastructure-level software (CUDA, NCCL, drivers, clusterware)?	The scope of Tier 3 scientific software support is comprehensive; the Government expects support for both infrastructure-level software and ML framework-level software.
9	Sec L, Go/No-Go #3; Tech Spec 5.1	(a) Does the NVIDIA RTX PRO 6000 Blackwell (or similar PCIe Blackwell workstation / pro-grade GPUs) satisfy the "equivalent or better" language in Go/No-Go Requirement 3, or is that requirement limited to SXM-form-factor datacenter GPUs with NVLink / NVSwitch fabric?	(a) No, the NVIDIA RTX PRO 6000 Blackwell (or similar PCIe workstation/pro-grade GPUs) does not satisfy the 'equivalent or better' language in Go/No-Go Requirement 3. The minimum baseline requirement is the NVIDIA HGX H100 80GB 700W SXM5

		(b) If not, would it be acceptable as part of the 100 GPU reserve capacity?	8X GPUs. The HGX H100's high-bandwidth fabric (NVSwitch) and memory architecture (HBM3) cannot be replicated in the RTX 6000 or similar PCIe form factors. (b) No, it would not be acceptable as part of the reserve capacity. As clarified in a recent amendment to the solicitation, the 100 additional reserve GPUs must meet the stated minimum baseline (H100, equivalent, or better).
10	Sec L, Go/No-Go #2; Tech Spec 5.3	A quoter may operate the same platform across two boundaries: a FedRAMP High / IL-5 boundary in GovCloud (accredited, with a full SSP maintained in eMASS) and a FISMA Low boundary in commercial cloud. Because commercial GPU providers do not currently offer an IL-5 IaaS package, a control plane connected to a commercial GPU provider operates on the FISMA Low boundary. (a) Does the Government accept evidence that a FISMA Low control-plane deployment uses the same platform, policies, procedures, base images, and cleared containers as the quoter's accredited FedRAMP High / IL-5 boundary — supported by the	(a) No, the Government will not evaluate eMASS security packages or control-mapping crosswalks as substitute evidence for compliance. Furthermore, this solicitation does not have Controlled Unclassified Information (CUI) requirements at this time, so a FedRAMP High or IL-5 boundary is not required. As stated in the solicitation, the required baseline is CMMC Level 1 (Self) alongside any other certifications explicitly named in the Go/No-Go criteria. Furthermore, the SOC 2 Type II certification needs to be demonstrated.

		full eMASS security package and a FedRAMP High to SOC 2 TSC control-mapping crosswalk — as satisfying Go/No-Go Requirement 2 at the control-plane layer? (b) At the GPU infrastructure layer, does the Government accept a current ISO 27001 certification (typically alongside ISO 9001, 14001, and 45001) in lieu of SOC 2 Type II?	(b) No. The Government will not accept an ISO 27001 certification in lieu of a SOC 2 Type II certification. Quoters must provide the required SOC 2 Type II certification to satisfy the Go/No-Go requirement.
11	Sec L, Factor 1; Sec M, Factor 1	If a quoter proposed a US-hosted control plane federating GPU compute hosted at an allied NATO-member data center — with all user authentication, SSH, and workflow traffic to the GPU cluster proxied through the US-hosted control plane (the US platform is the sole user ingress/egress path; users do not connect directly to the non-US data center) — how would the Government weight that under Factor 1, relative to a fully offshore-hosted offer? Does the US-hosted control plane materially mitigate the Factor 1 US-preference weighting applied to the GPU compute layer?	The Government will evaluate all proposals in accordance with the criteria stated in the RFQ. As noted in Section L, under Factor 1 (Technical Approach), 'the Government's evaluation will give preference to solutions hosted entirely within the United States and its territories.' The degree to which any proposed solution, including a hybrid architecture, meets the Government's objectives will be evaluated as part of the tradeoff process described in Section M, where the Government will assess the relative strengths, weaknesses, and risks of each offeror's technical approach.

12	N/A	The solicitation specifies CMMC Level 1 (Self), which is consistent with unclassified / Federal Contract Information (FCI) workloads. Does DARPA anticipate any NODES data being CUI (32 CFR 2002), ITAR-controlled (22 CFR 120-130), or EAR-controlled (15 CFR 730-774) during the period of performance? If yes, the eligible GPU infrastructure set would be limited to FedRAMP High / IL-5 hyperscale providers (AWS GovCloud, Azure Government, Google Cloud for Government, Oracle US Government Cloud); commercial GPU providers would be ineligible. Confirmation before 5 May 2026 would allow quoters to re-scope provider shortlists and pricing.	No. DARPA does not anticipate processing any CUI, ITAR, or EAR data during the period of performance for this anticipated award resulting from this solicitation. Quoters should note, however, that future continuations of this work under the broader NODES program may have CUI requirements, as further specified in Attachment 1- Technical Specifications Documents and Attachment 2- Instructions to Quoters.
13	N/A	(a) Does the Government permit a quote to name a primary GPU infrastructure provider while preserving the contractor's ability to migrate capacity to a secondary qualified provider during the period of performance (for capacity shift, pricing change, or newer hardware availability)? (b) If so, is the preferred mechanism advance notification to the Contracting Officer, or a formal contract modification?	(a) The Government permits and encourages quoters to propose their best technical solution. All quotes will be evaluated in accordance with Section M of the solicitation. If an offeror intends to use a primary GPU infrastructure provider while preserving the ability to migrate capacity to a secondary qualified provider, that proposed technical solution, including the circumstances, processes, and potential pricing impacts of such a migration must be described in sufficient detail in the quote to allow for a thorough

			evaluation.(b) Any material changes after award will require a formal, bilateral contract modification.
14	Section L, (b) Quotation Format	(a) Do appendices or supporting attachments to Volume I (manufacturer datasheets, benchmark results, architecture diagrams, reference letters) count against the Volume I 10-page limit, or may they be provided as separately evaluated supplements? (b) Is there a preferred location for benchmark data and manufacturer specifications supporting the "equivalent or better" or "advanced technology" language?	(a) These items will count against the Volume I 10 page limit unless they serve the purpose of demonstrating technology superiority as in (b) below. (b) Supporting evidence for higher performance of proposed technology can be excluded from the Volume I page limitation as long as it's labeled appropriately ("benchmark data/manufacturer specification to support equivalency or superiority").
15	Section M, Factor 3 (Price)	Is there a preferred contract mechanism for proposing a GPU hardware refresh at the option-period exercise (e.g., H100 baseline to Blackwell) as a priced option within the initial quote rather than via a subsequent modification? Would such a priced option compete under Section M Price evaluation, or be evaluated as a separate optional add-on?	To support simple, consumption-based billing, the Government strongly prefers all costs, including hardware refreshes, be rolled into the all-inclusive per-GPU-hour rate. As outlined in Solicitation Amendment 1, quoters may utilize the newly added CLIN 0003 for ancillary services that cannot reasonably be included in the hourly rate. Any costs proposed under CLIN 0003 will be

		We appreciate the Government's consideration of these questions and will incorporate the responses into the 5 May 2026 quote submission.	incorporated into the Total Evaluated Price under Section M.
--	--	--	--