

Market Research Questionnaire

The purpose of this document is for market research. This document does not commit or obligate the Government in any manner. Please complete all fields and provide documentation to support your responses below. Failure to respond fully to this questionnaire will render your response incomplete and your company's response will not be considered a viable source. Fill in the form completely and send to Parvinder.s.brar@irs.gov no later than May 5, 2026, 2PM PST.

1a. Service Title: **Multimodal AI Investigative Platform**

1b. website _____

Company Information:

2. Company legal name: _____

3. POC: _____ Email: _____ Phone: _____

4. UEI#: _____

5. Cage Code: _____

6. System for Award Management (SAM) Registered and active?

7. Any SAM exclusions: _____

7a. NAICS: _____ 7b. Business Size: _____

7c. Socio-Economic Category (8(a), HUBZone, WOSB, SDVOSB etc.): _____

8. GSA contract# (if applicable): _____

Past Experience: (Note: past experience is not required to be considered.)

10. List company's past experience providing this service with the Federal government?

Product	Contract #	Agency	Govt POC Name	Phone	Email
_____	_____	_____	_____	_____	_____
_____	_____	_____	_____	_____	_____
_____	_____	_____	_____	_____	_____
_____	_____	_____	_____	_____	_____

Note: Additional items can be added on continuation page

Service Information:

11. Price per unit: _____

12. Desired line item structure, e.g. what is the unit of measure for invoices? (Per hour, per job, per month, per square foot)

13. Please provide pricing information for previous projects:

Size Status Information:

14. The primary purpose of this acquisition is to provide **Multimodal AI Investigative Platform**:

14a. Of the service identified in the SOW/PWS, which one(s) will your company perform as the prime contractor using its own personnel without subcontracting?

14b. With respect to the service actions identified in the SOW/PWS, which action(s) do you intend on subcontracting or using other than your company's own personnel?

14c. Of those services that will be performed by subcontractor(s), approximately what percentage will they perform for each of the actions identified?

14d. What is the size status of the intended subcontractor(s)? _____

14e. If you intend on subcontracting, please detail how you will be in compliance with the Limitations on Subcontracting.

Technical Questions:

1. Evidentiary Traceability & Legal Defensibility

Describe how your solution establishes and maintains deterministic, immutable linkage between all derived data elements and their originating source artifacts.

Your response shall include:

- A detailed explanation of how source-to-output traceability is preserved at the data element level (e.g., page, timestamp, extraction method).
- Description of how version control and corrections are handled without breaking lineage.
- Sample machine-readable provenance output (e.g., JSON, XML, or structured log).
- Explanation of how this capability operates identically in air-gapped environments.
- Identification of any probabilistic or inferred linkage mechanisms, if applicable, and how they are mitigated for evidentiary use.

2. Deployment Architecture & Government Operational Independence

Provide a detailed technical description of how your solution operates fully within Government-controlled, on-premises and air-gapped environments without reliance on vendor-hosted services.

Your response shall include:

- Architecture diagram showing deployment in:
 - Fully disconnected (air-gapped)
 - Restricted network
 - Connected environments

- Explanation of how AI/ML and LLM capabilities are deployed locally, including model hosting and update mechanisms.
- Description of how the Government can:
 - Configure, maintain, and extend the platform without vendor-embedded personnel
 - Scale compute resources (including GPU acceleration) without architectural redesign
- Identification of any external dependencies, telemetry, or call-home features (must be explicitly disclosed).

3. Enterprise Throughput, Concurrency & Performance Validation

Demonstrate how your solution meets or exceeds the Government's enterprise-scale workload requirements.

Your response shall include:

- Documented performance metrics for:
 - Concurrent users (target ~2,500 users)
 - High-volume ingestion (financial documents, multimedia >4GB)
 - Token processing capacity ($\geq$ 1M tokens)
- Description of how the platform prevents:
 - License contention
 - Processing throttling
 - Operational backlog
- Real-world or production benchmark data (not theoretical estimates).
- Explanation of scaling model (compute vs licensing vs throughput).
- Identification of any hard limits (users, files, processing, tokens, etc.).

4. Salient Characteristics Compliance Matrix (Required Chart Completion)

Responders shall complete all fields in the matrix below.

Responses shall be specific, technical, and verifiable. Narrative-only responses are not sufficient.

- Compliance must be marked as:
Y = Fully Compliant | P = Partial | N = Not Compliant
- "Partial" responses must include mitigation approach, timeline, and associated risk
- "Custom" indicates capability requires non-native development, integration, or modification
- Failure to complete all fields may result in an unfavorable rating

Matrix

#	Salient Characteristic Area	Requirement	Compliance (Y/P/N)	Description of How Requirement is Met	Native / Configured / Custom
1.1	Deployment & Infrastructure	Operates fully on-prem, air-gapped, no SaaS dependency			
1.2	Deployment & Infrastructure	Functional parity across connected and disconnected environments			
1.3	Deployment & Infrastructure	Operable and extensible without vendor-embedded staff			
1.4	Deployment & Infrastructure	Supports CI-2 HCI and composable GPU infrastructure (Base Period integration)			
1.5	Deployment & Infrastructure	Scales via Government-controlled compute without external data movement			
2.1	Performance & Throughput	Demonstrates sustained enterprise-scale processing capacity			
2.2	Performance & Throughput	Supports high-volume intake without backlog			
2.3	Performance & Throughput	Maintains reliability under peak workload conditions			
2.4	Performance & Throughput	Supports enterprise concurrent users (~2,500 users)			
2.5	Performance & Throughput	Scalable ingestion aligned to annual intake volumes			
2.6	Performance & Throughput	Supports ≥1M tokens			
3.1	Data Ownership & Portability	Government retains full ownership of all data and outputs			
3.2	Data Ownership & Portability	Outputs exportable in non-proprietary formats			
3.3	Data Ownership & Portability	No proprietary ontologies restricting use or migration			
3.4	Data Ownership & Portability	No vendor lock-in or egress barriers			
4.1	Multimodal AI	Processes structured, semi-structured, and unstructured data			
4.2	Multimodal AI	Performs OCR, handwriting recognition, extraction, transcription, translation locally			
4.3	Multimodal AI	AI/LLM operates fully within Government-controlled environments			
4.4	Multimodal AI	Automated classification and workflow routing			
4.5	Multimodal AI	Human-in-the-loop validation,			

		override, and audit logging			
4.6	Multimodal AI	Full transparency and traceability of AI outputs			
5.1	Financial Forensics	Processes financial artifacts (bank statements, checks, etc.)			
5.2	Financial Forensics	Supports continuously updated financial format libraries			
5.3	Financial Forensics	Automated categorization, classification, reconciliation, anomaly detection			
5.4	Financial Forensics	Outputs suitable for evidentiary use (FRE 901 & 702)			
6.1	Large File & Media	Processes large media files without segmentation			
6.2	Large File & Media	Preserves timestamps, metadata, and source continuity			
6.3	Large File & Media	Parallel processing without workflow throttling			
6.4	Large File & Media	Multilingual transcription/translation in Government environment			
6.5	Large File & Media	Supports Chinese dialects (e.g., Cantonese, Shanghainese, Jin)			
6.6	Large File & Media	Maintains linkage between original and translated content			
6.7	Large File & Media	Supports $\geq 1M$ tokens			
7.1	Traceability & Provenance	Deterministic linkage between data and source artifacts			
7.2	Traceability & Provenance	Data elements include source ID, page/timestamp, extraction method			
7.3	Traceability & Provenance	Immutable lineage with version-controlled corrections			
7.4	Traceability & Provenance	Direct navigation from output to source evidence			
7.5	Traceability & Provenance	Machine-readable provenance logs exportable			
7.6	Traceability & Provenance	Operates identically in offline/air-gapped environments			
8.1	Governance & Explainability	Supports automated triage and decision-support workflows			
8.2	Governance & Explainability	Outputs are explainable and attributable to source artifacts			
8.3	Governance & Explainability	Mandatory human-in-the-loop review and override			
8.4	Governance & Explainability	Monitors performance drift, bias, degradation			

8.5	Governance & Explainability	Logs system changes, retraining, and user overrides			
9.1	Licensing & Sustainability	Supports enterprise licensing (~2,500 users)			
9.2	Licensing & Sustainability	No throttling during surge investigative operations			
9.3	Licensing & Sustainability	No per-use/per-transaction licensing constraints			
9.4	Licensing & Sustainability	Scales compute without duplicative licensing			
9.5	Licensing & Sustainability	Transparent enterprise pricing model			

Failure to demonstrate full compliance (Y) with any requirement that impacts evidentiary traceability, deployment independence, or enterprise scalability may result in an unfavorable rating.

Responses marked “Partial” will be reviewed for risk, timeline, and operational impact, and may be deemed unfavorable if they introduce material risk to Government investigative operations.